
CS5785 Homework 3

The homework is generally split into programming exercises and written exercises.

This homework is due on **November 10, 2015 at 11:59 PM EST**. Upload your homework to [CMS](#).

Please upload as a single .zip file. A complete submission should include:

- A write-up as a single .pdf file;
- Source code and data files for all of your experiments (AND figures) in .py files if you use Python or .ipynb files if you use the IPython Notebook. If you use some other language, include all build scripts necessary to build and run your project.

The write-up should be in **professional lab report format**. It should contain a general summary of what you did, how well your solution works, any insights you found, etc. On the cover page, include the class name, homework number, and team member names. You are responsible for submitting clear, organized answers to the questions.

Please include all relevant information for a question, including text response, equations, figures, graphs, output, etc. If you include graphs, be sure to include the source code that generated them.

Please pay attention to the discussion board for relevant information regarding updates, tips, and policy changes. You are encouraged (but not required) to work in groups of 2.

IF YOU NEED HELP

There are several strategies available to you.

- If you ever get stuck, the best way is to ask your teammates on Piazza.¹ That way, your solutions will be available to the other students in the class.
- Your instructor and TAs will offer office hours,² which are a great way to get some one-on-one help.

¹<http://piazza.com/cornell/fall2015/cs5785>

²<http://cs5785-cornell-tech.github.io/contact.html>

PROGRAMMING EXERCISES

1. Sentiment analysis of online reviews.

In this assignment you will use several machine learning techniques from the class to identify and extract subjective information in online reviews. Specifically, the task for this assignment is **sentiment analysis**. According to *Wikipedia*, sentiment analysis aims to determine the attitude of a speaker or a writer with respect to some topic or the overall contextual polarity of a document. It has been shown that people's attitudes are largely manifested in the language they adopt. This assignment will walk you through the mystery and help you better understand our posts online!

Important: Use your own implementations for this assignment. Any K-means, bag-of-words or PCA implementation is **NOT** allowed.

- (a) Download [Sentiment Labelled Sentences Data Set](#). There are three data files under the root folder. *yelp_labelled.txt*, *amazon_cells_labelled.txt* and *imdb_labelled.txt*. Parse each file with the specifications in *readme.txt*. Are the labels balanced? If not, what's the ratio between the two labels? Explain how you process these files.
- (b) **Pick your preprocessing strategy.** Since these sentences are online reviews, they may contain significant amounts of noise and garbage. You **may or may not** want to do one or all of the following. Explain the reasons for each of your decision (**why or why not**).
 - Lowercase all of the words.
 - Lemmatization of all the words.
 - Strip punctuation.
 - Strip the stop words, e.g., "the", "and", "or".
 - Something else? Tell us about it.
- (c) **Split training and testing set.** In this assignment, for each file, please use the first 400 instances for each label as the *training set* and the remaining 100 instances as *testing set*. In total, there are 2400 reviews for training and 600 reviews for testing.
- (d) **Bag of Words model.** Extract features and then represent each review using bag of words model, i.e., every word in the review becomes its own element in a feature vector. In order to do this, first, make one pass through all the reviews in the *training set* (**Explain why** we can't use *testing set* at this point) and build a dictionary of unique words. Then, make another pass through the review in both the *training set* and *testing set* and count up the occurrences of each word in your dictionary. The i th element of a review's feature vector is the number of occurrences of the i th dictionary word in the review. Implement the bag of words model and report feature vectors of any two reviews in the training set.
- (e) **Pick your postprocessing strategy.** Since the vast majority of English words will not appear in most of the reviews, most of the feature vector elements will be 0. This suggests that we need a postprocessing or normalization strategy that combats the huge variance of the elements in the feature vector. You may want to use one of the following strategies. Whatever choices you make, explain **why** you made the decision.
 - *log-normalization*. For each element of the feature vector x , transform it into $f(x) = \log(x + 1)$.
 - *l1 normalization*. Normalize the $l1$ norm of the feature vector, $\hat{\mathbf{x}} = \frac{\mathbf{x}}{|\mathbf{x}|}$.

- *l2 normalization*. Normalize the *l2* norm of the feature vector, $\hat{\mathbf{x}} = \frac{\mathbf{x}}{\|\mathbf{x}\|}$.
 - *Standardize* the data by subtracting the mean and dividing by the variance.
- (f) **Training set clustering.** For the feature vectors you computed from (e). Implement K-means algorithm to divide the training set into 2 clusters (i.e., $K = 2$). Report the centers of both clusters. Inspect the clustering results and the labels associated the each instance, and evaluate the performance of your K-means algorithm and bag of words feature representation.
- (g) **Sentiment prediction.** Train a logistic regression model (*you can use existing packages here*) on the training set and test on the testing set. Report the classification accuracy and confusion matrix. Inspecting the weight vector of the logistic regression, what are the words that play the most important roles in deciding the sentiment of the reviews?
- (h) **N-gram model.** Similar to the bag of words model, but now you build up a dictionary of n-grams, which are contiguous sequences of words. For example, "Alice fell down the rabbit hole" would then map to the 2-grams sequence: ["Alice fell", "fell down", "down the", "the rabbit", "rabbit hole"], and all five of those symbols would be members of the n-gram dictionary. Try $n = 2$, repeat (d)-(g) and report your results.
- (i) **PCA for bag of words model.** The features in the bag of words model have large redundancy. Implement PCA to reduce the dimension of features calculated in (e) to 10, 50 and 100 respectively. Using these lower-dimensional feature vectors and repeat (f), (g). Report corresponding clustering and classification results. (**Note:** You should implement PCA yourself, but you can use `numpy.svd` or some other SVD package. Feel free to double-check your PCA implementation against an existing one)
- (j) **Algorithms comparison and analysis.** According to the above results, compare the performances of *bag of words*, *2-gram* and *PCA for bag of words*. Which method performs best in the prediction task and why? What do you learn about the language that people use in on-line reviews (e.g., expressions that will make the posts positive/negative)? *Hint:* Inspect the clustering results and the weights learned from logistic regression.

2. EM algorithm and implementation

- (a) The parameters of Gaussian Mixture Model (GMM) can be estimated via the EM algorithm. Show that the [alternating algorithm](#) for K-means is a special case of the EM algorithm and show the corresponding objective functions for E-step and M-step.
- (b) Download the [Old Faithful Geyser Dataset](#). The data file contains 272 observations of (*eruption time*, *waiting time*). Treat each entry as a 2 dimensional feature vector. Parse and plot all data points on 2-D plane.
- (c) Implement a bimodal GMM model to fit all data points using EM algorithm. Explain the reasoning behind your termination criteria. For this problem, we assume the covariance matrix is spherical (i.e., it has the form of $\sigma^2 I$) and you can randomly initialize Gaussian parameters. For evaluation purposes, please submit the following figures:
- Plot the trajectories of two mean vectors in 2 dimensions (i.e., coordinates vs. iteration).
 - Run your program for 50 times with different initial parameter guesses. Show the distribution of the total number of iterations needed for algorithm to converge.

(d) Repeat the task in (c) but with the initial guesses of the parameters generated from the following process:

- Run a K-means algorithm over all the data points with $K = 2$ and label each point with one of the two clusters.
- Estimate the first guess of the mean and covariance matrices using maximum likelihood over the labeled data points.

Compare the algorithm performances of (c) and (d).

WRITTEN EXERCISES

You can find links to the textbooks for our class on [the course website](#).

Submit the answers to these questions along with your writeup as a single .pdf file. We do recommend you to type your solutions using LaTeX or other text editors, since hand-written solutions are often hard to read. If you handwrite them, please be legible!

1. HTF Exercise 14.2 (Gaussian mixture model and EM algorithm)
2. HTF Exercise 14.8 (Procrustes problem.)
3. HTF Exercise 14.11 (Multidimensional Scaling.)